

Credit Card Default

Aymen Rumi

5/8/2020

Overview

We will analyze and build a model for Credit Card Default data, with response being binary variable “default” and predictors being binary variable “student” and continuous variables balance & income

Plotting & Visualization

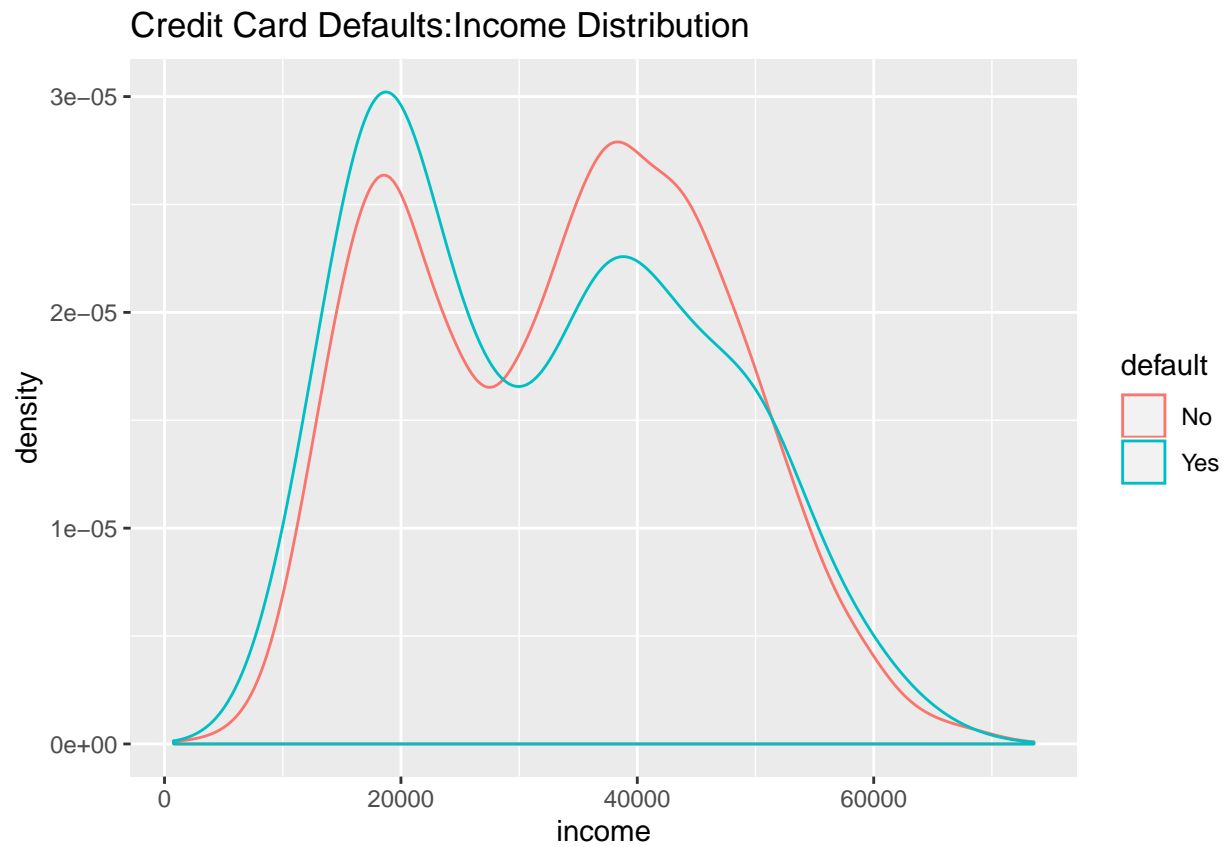
We explore the relationships between our response and predictors with visualizations

```
ggplot(data=Default,aes(x=income,y=balance,color=default))+geom_point()+stat_ellipse()+ggtitle("Credit Card Defaults: Scatterplot")
```

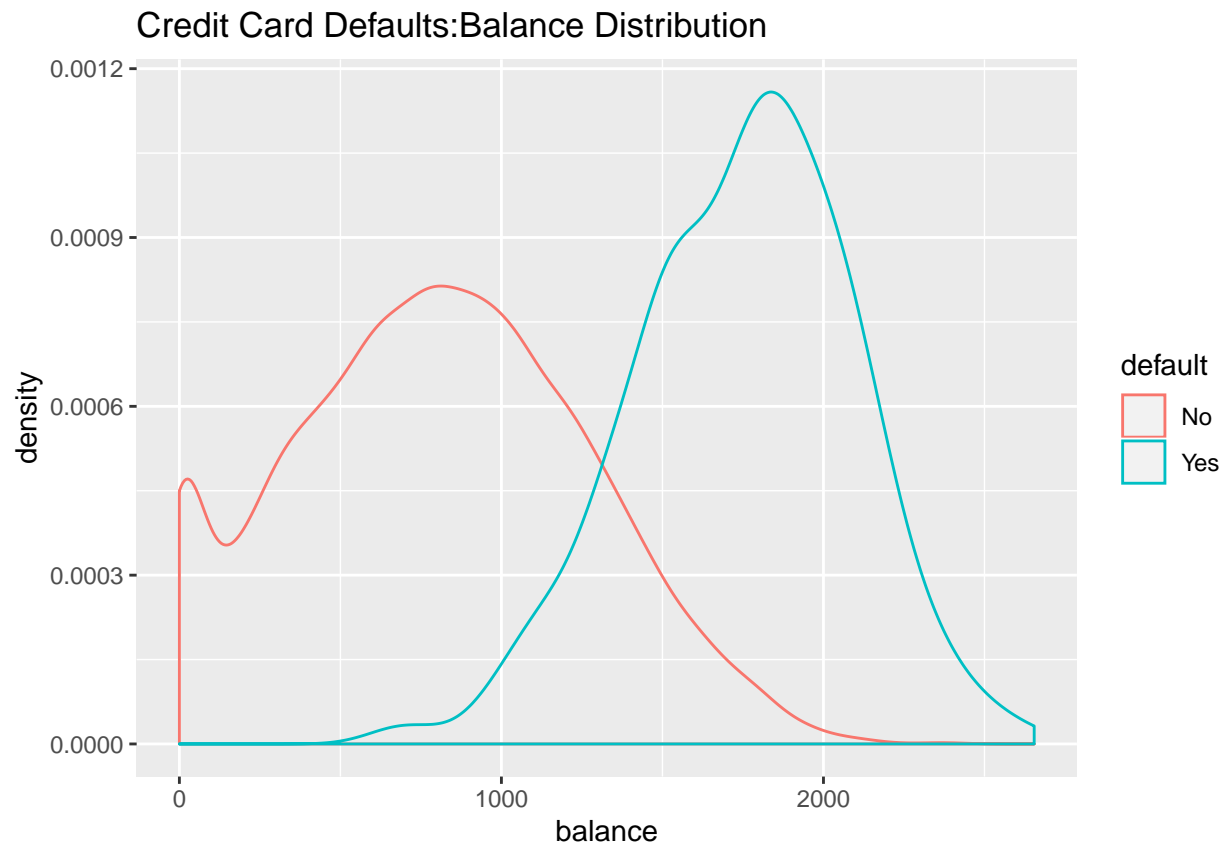
Credit Card Defaults: Scatterplot



```
ggplot(data=Default,aes(x=income,color=default))+geom_density()+ggtitle("Credit Card Defaults: Income Distribution")
```



```
ggplot(data=Default,aes(x=balance,color=default))+geom_density()+ggtitle("Credit Card Defaults:Balance I
```




```
ggplot(data=Default,aes(x=income,y=balance,color=student))+geom_point()+ggtitle("Students: Scatterplot")
```

Students: Scatterplot




```
ggplot(data=Default, aes(x=income,y=balance, fill=student)) +ggtitle("Students: Boxplot")+  
geom_boxplot()
```

Students: Boxplot




```
ggplot(data=subset(Default,default=="Yes"),aes(x=income,y=balance,color=student))+geom_point()+ggtitle(
```

Students(Defaulted): Scatterplot




```
ggplot(data=subset(Default,default=="Yes"), aes(x=income,y=balance, fill=student)) +  
  geom_boxplot()+ggtitle("Students(Defaulted): Boxplot")
```

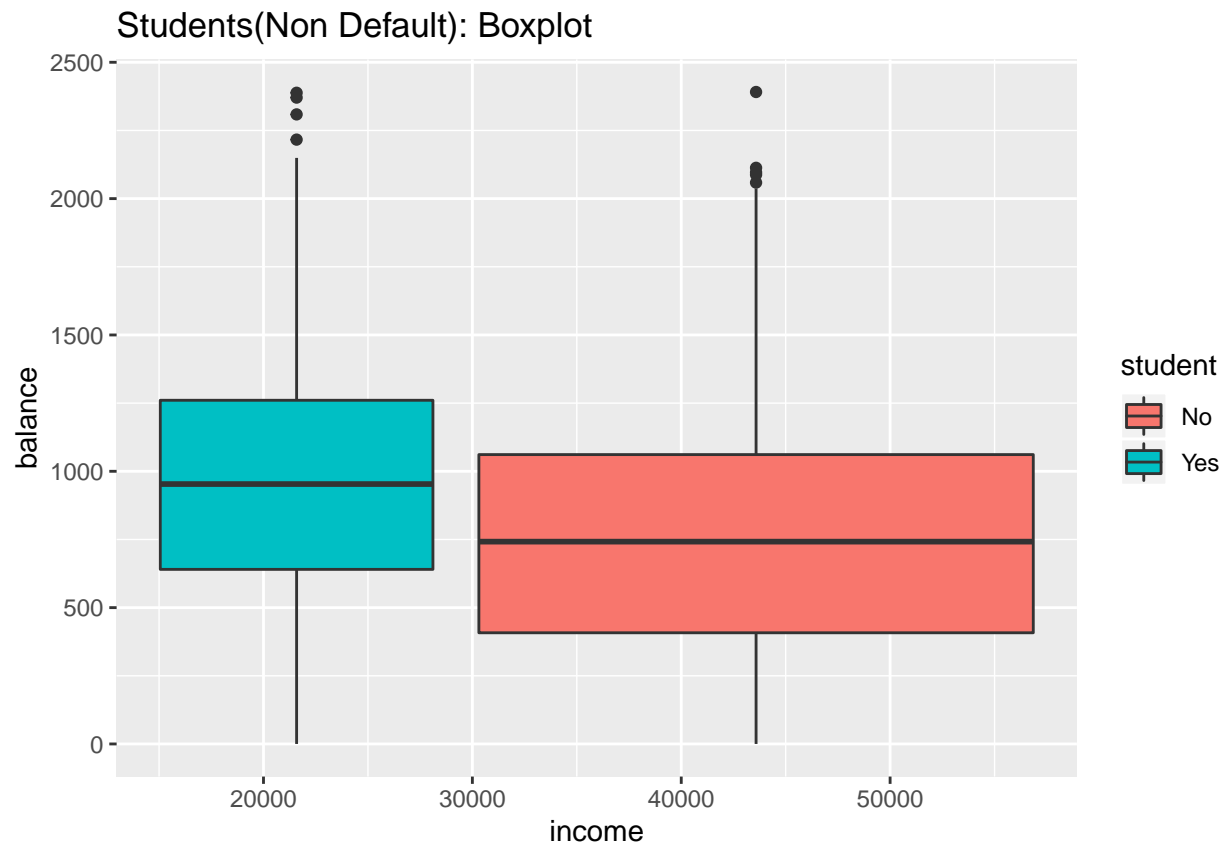
Students(Defaulted): Boxplot



```
ggplot(data=subset(Default,default=="No"),aes(x=income,y=balance,color=student))+geom_point()+ggtitle("Students(Defaulted): Boxplot")
```



```
ggplot(data=subset(Default,default=="No"), aes(x=income,y=balance, fill=student)) +  
  geom_boxplot()+ggtitle("Students(Non Default): Boxplot")
```

Logistic Regression Model Training

```
attach(Default)
set.seed(1)

glm.fit=glm(default~balance+income,family=binomial,data=Default)

glm.probs=predict(glm.fit,type="response")

glm.pred=rep("No",length(Default$default))
glm.pred[glm.probs>0.5]="Yes"

table(glm.pred,default)
```

```
##          default
## glm.pred  No  Yes
##      No  9629  225
##      Yes   38  108
```

```
mean(glm.pred==default)
```

```
## [1] 0.9737
```

```
mean(glm.pred!=default)
```

```
## [1] 0.0263
```

K-Fold Cross Validation

```
cv.err=cv.glm(Default,glm.fit,K=5)  
cv.err$delta[1]
```

```
## [1] 0.02149045
```

```
1-cv.err$delta[1]
```

```
## [1] 0.9785095
```

```
cv.err=cv.glm(Default,glm.fit,K=10)  
cv.err$delta[1]
```

```
## [1] 0.02147942
```

```
1-cv.err$delta[1]
```

```
## [1] 0.9785206
```

```
cv.err=cv.glm(Default,glm.fit,K=100)  
cv.err$delta[1]
```

```
## [1] 0.0214586
```

```
1-cv.err$delta[1]
```

```
## [1] 0.9785414
```

Bootstrap

```
classification_estimate<-function(data,index)  
{  
  response<-data$default[index]  
  x1<-data$balance[index]  
  x2<-data$income[index]  
  
  glm.fit=glm(response~x1+x2,famil=binomial)  
  
  glm.probs=predict(glm.fit,type="response")  
}
```

```

glm.pred=rep("No",length(Default$default))
glm.pred[glm.probs>0.5]="Yes"

table(glm.pred,default)

return (mean(glm.pred==default))
}

boot(Default,statistic = classification_estimate,R=1000)

##
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
##
## Call:
## boot(data = Default, statistic = classification_estimate, R = 1000)
##
##
## Bootstrap Statistics :
##      original      bias    std. error
## t1*    0.9737 -0.0205257 0.001481616

```