

Model Selection Methods

Aymen Rumi

5/20/2020

Overview

We will run regression analysis on the College. The goal in this report is to try out different model selection techniques for regression analysis such as variable selection, regularization & dimensionality reduction techniques, in aims to pick the best predictive model for our dataset(lowest error rate on testing data)

College Dataset

We want to predict the number of applicants using all the variables available

Standard Regression Analysis

```
summary(lm(data=College,Apps~.))
```

Call:

```
lm(formula = Apps ~ ., data = College)
```

Residuals:

Min	1Q	Median	3Q	Max
-4908.8	-430.2	-29.5	322.3	7852.5

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-445.08413	408.32855	-1.090	0.276053	
PrivateYes	-494.14897	137.81191	-3.586	0.000358	***
Accept	1.58581	0.04074	38.924	< 2e-16	***
Enroll	-0.88069	0.18596	-4.736	2.60e-06	***
Top10perc	49.92628	5.57824	8.950	< 2e-16	***
Top25perc	-14.23448	4.47914	-3.178	0.001543	**
F.Undergrad	0.05739	0.03271	1.754	0.079785	.
P.Undergrad	0.04445	0.03214	1.383	0.167114	
Outstate	-0.08587	0.01906	-4.506	7.64e-06	***
Room.Board	0.15103	0.04829	3.127	0.001832	**
Books	0.02090	0.23841	0.088	0.930175	
Personal	0.03110	0.06308	0.493	0.622060	
PhD	-8.67850	4.63814	-1.871	0.061714	.
Terminal	-3.33066	5.09494	-0.654	0.513492	
S.F.Ratio	15.38961	13.00622	1.183	0.237081	
perc.alumni	0.17867	4.10230	0.044	0.965273	
Expend	0.07790	0.01235	6.308	4.79e-10	***
Grad.Rate	8.66763	2.94893	2.939	0.003390	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1041 on 759 degrees of freedom
Multiple R-squared: 0.9292, Adjusted R-squared: 0.9276
F-statistic: 585.9 on 17 and 759 DF, p-value: < 2.2e-16

Observations

From the regression function summary, 10 variables seem to have a low p-value indicating significant contribution to predicting the response, these variables are: PrivateYes, Accept,Enroll, Top10perc, Top25perc,Outstate,PhD,Room.Board, Expend,Grad.Rate

Selecting Best Subset of Predictors

We will perform forward stepwise variable selection, & plot graphs for statistics such as Cp,BIC,RSS, & Rsquared as the # of variables increases

```
College<-College

regfit.full=regsubsets(Apps~.,data=College,nvmax=17,method="forward")
reg.summary=summary(regfit.full)

reg.summary
```

Subset selection object

Call: regsubsets.formula(Apps ~ ., data = College, nvmax = 17, method = "forward")

17 Variables (and intercept)

	Forced in	Forced out
PrivateYes	FALSE	FALSE
Accept	FALSE	FALSE
Enroll	FALSE	FALSE
Top10perc	FALSE	FALSE
Top25perc	FALSE	FALSE
F.Undergrad	FALSE	FALSE
P.Undergrad	FALSE	FALSE
Outstate	FALSE	FALSE
Room.Board	FALSE	FALSE
Books	FALSE	FALSE
Personal	FALSE	FALSE
PhD	FALSE	FALSE
Terminal	FALSE	FALSE
S.F.Ratio	FALSE	FALSE
perc.alumni	FALSE	FALSE
Expend	FALSE	FALSE
Grad.Rate	FALSE	FALSE

1 subsets of each size up to 17

Selection Algorithm: forward

		PrivateYes	Accept	Enroll	Top10perc	Top25perc	F.Undergrad
1	(1)	" "	"*	" "	" "	" "	" "
2	(1)	" "	"*	" "	"*	" "	" "
3	(1)	" "	"*	" "	"*	" "	" "
4	(1)	" "	"*	" "	"*	" "	" "
5	(1)	" "	"*	"*	"*	" "	" "

6	(1)	" "	"*"	"*"	"*"	" "	" "
7	(1)	" "	"*"	"*"	"*"	"*"	" "
8	(1)	"*"	"*"	"*"	"*"	"*"	" "
9	(1)	"*"	"*"	"*"	"*"	"*"	" "
10	(1)	"*"	"*"	"*"	"*"	"*"	" "
11	(1)	"*"	"*"	"*"	"*"	"*"	"*"
12	(1)	"*"	"*"	"*"	"*"	"*"	"*"
13	(1)	"*"	"*"	"*"	"*"	"*"	"*"
14	(1)	"*"	"*"	"*"	"*"	"*"	"*"
15	(1)	"*"	"*"	"*"	"*"	"*"	"*"
16	(1)	"*"	"*"	"*"	"*"	"*"	"*"
17	(1)	"*"	"*"	"*"	"*"	"*"	"*"

		P.Undergrad	Outstate	Room.Board	Books	Personal	PhD	Terminal
1	(1)	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	"*"	" "	" "	" "	" "	" "
5	(1)	" "	"*"	" "	" "	" "	" "	" "
6	(1)	" "	"*"	"*"	" "	" "	" "	" "
7	(1)	" "	"*"	"*"	" "	" "	" "	" "
8	(1)	" "	"*"	"*"	" "	" "	" "	" "
9	(1)	" "	"*"	"*"	" "	" "	"*"	" "
10	(1)	" "	"*"	"*"	" "	" "	"*"	" "
11	(1)	" "	"*"	"*"	" "	" "	"*"	" "
12	(1)	"*"	"*"	"*"	" "	" "	"*"	" "
13	(1)	"*"	"*"	"*"	" "	" "	"*"	" "
14	(1)	"*"	"*"	"*"	" "	" "	"*"	"*"
15	(1)	"*"	"*"	"*"	" "	"*"	"*"	"*"
16	(1)	"*"	"*"	"*"	"*"	"*"	"*"	"*"
17	(1)	"*"	"*"	"*"	"*"	"*"	"*"	"*"

		S.F.Ratio	perc.alumni	Expend	Grad.Rate
1	(1)	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "
3	(1)	" "	" "	"*"	" "
4	(1)	" "	" "	"*"	" "
5	(1)	" "	" "	"*"	" "
6	(1)	" "	" "	"*"	" "
7	(1)	" "	" "	"*"	" "
8	(1)	" "	" "	"*"	" "
9	(1)	" "	" "	"*"	" "
10	(1)	" "	" "	"*"	"*"
11	(1)	" "	" "	"*"	"*"
12	(1)	" "	" "	"*"	"*"
13	(1)	"*"	" "	"*"	"*"
14	(1)	"*"	" "	"*"	"*"
15	(1)	"*"	" "	"*"	"*"
16	(1)	"*"	" "	"*"	"*"
17	(1)	"*"	"*"	"*"	"*"

```
par(mfrow=c(2,2))
```

```
which.max(reg.summary$adjr2)
```

```
[1] 13
```

```
plot(reg.summary$adjr2,xlab="Number of Variables",
      ylab="Adjusted RSq",type="l")
points(which.max(reg.summary$adjr2),reg.summary$adjr2[which.max(reg.summary$adjr2)], col="red",cex=2,pch=20)

which.min(reg.summary$rss)
```

[1] 17

```
plot(reg.summary$adjr2,xlab="Number of Variables",
      ylab="RSS",type="l")
points(which.min(reg.summary$rss),reg.summary$adjr2[which.min(reg.summary$rss)], col="red",cex=2,pch=20)

which.min(reg.summary$cp)
```





[1] 12

```
plot(reg.summary$cp,xlab="Number of Variables",
      ylab="Cp",type="l")
points(which.min(reg.summary$cp),reg.summary$cp[which.min(reg.summary$cp)], col="red",cex=2,pch=20)

which.min(reg.summary$bic)
```

[1] 10

```
plot(reg.summary$bic,xlab="Number of Variables",
      ylab="BIC",type="l")
points(which.min(reg.summary$bic),reg.summary$bic[which.min(reg.summary$bic)], col="red",cex=2,pch=20)
```



Observations

As we can see there is still no clear evidence of the optimal number of variables to use but BIC had the same result as the regression model with 10 variables: PrivateYes, Accept, Enroll, Top10perc, Top25perc, Outstate, Room.Board, PhD, Expend, Grad.Rate.

The goal of choosing the best subset of variables is to optimize predictive accuracy on future unseen dataset, so let's analyze performance of these subsets by calculating test error rate using K-Fold Cross Validation

Validation Set

```
set.seed(4)

# split the dataset

sample <- sample.split(College$Apps, SplitRatio = .70)

train <- subset(College, sample == TRUE)
test  <- subset(College, sample == FALSE)

regfit.fwd=regsubsets(Apps~.,data=train, nvmax=17,method="forward")

test.mat=model.matrix(Apps~.,data=test)

validation.errors=rep(NA,17)

for(i in 1:17){
  coefi=coef(regfit.fwd,id=i)
  pred=test.mat[,names(coefi)]%*%coefi
  validation.errors[i]=mean((test$Apps-pred)^2)
}

validation.errors

[1] 1820930 1583796 1705665 1688576 1651695 1657382 1870889 1922150
[9] 1952357 1939974 1870591 1852722 1854180 1853134 1850381 1849124
[17] 1855638
```

```
which.min(validation.errors)
```

```
[1] 2
```

```
coef(regfit.fwd,id=which.min(validation.errors))
```

```
(Intercept)      Accept      Top10perc
-824.439455      1.362166      38.437432
```

K-Fold Cross Validation

```
set.seed(1)
k=10
folds=sample(1:k,nrow(College),replace=TRUE)
cv.errors=matrix(NA,k,17, dimnames=list(NULL, paste(1:17)))

predict.regsubsets=function(object,newdata,id,...){
  form=as.formula(object$call[[2]])
  mat=model.matrix(form,newdata)
  coef=coef(object,id=id)
  xvars=names(coef)
  mat[,xvars] %*% coef
}

for(j in 1:k)
{
  best.fit=regsubsets(Apps~.,data=College[folds!=j,],nvmax=17)

  for (i in 1:17)
  {
    pred=predict.regsubsets(best.fit,College[folds==j,],id=i)
    cv.errors[j,i]=mean((College$Apps[folds==j]-pred)^2)
  }
}

mean.cv.errors=apply(cv.errors,2,mean)

which.min(mean.cv.errors)
```

4
4

Observations

If we plot minimum instance of test error we can see that different variables are chosen at different iterations, confirming a high level of overfitting to test data, best method to tackle this will be using regularization

L1-Lasso Regularization & L2-Ridge Regularization

```
x=model.matrix(Apps~.,College)[,-1]
y=College$Apps

# Ridge Regression

grid=10^seq(10,-2,length=100)

ridge.mod=glmnet(x,y,alpha=0,lambda=grid)
```

```

# Train & Test

train=sample(1:nrow(x),nrow(x)/2)

test=(-train)

y.test=y[test]

ridge.mod=glmnet(x[train,],y[train],alpha=0,lambda=grid,thresh=1e-12)

ridge.pred=predict(ridge.mod,s=4,newx=x[test,])

mean((ridge.pred-y.test)^2)

```

```
[1] 1162752
```


```

# cross validation

set.seed(1)

cv.out=cv.glmnet(x[train,],y[train],alpha=0)
plot(cv.out)

```



```

bestlam=cv.out$lambda.min

bestlam

```


```
[1] 390.979
```

```
ridge.pred=predict(ridge.mod,s=bestlam,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
[1] 1062563
```

```
# lasso regression
```

```
cv.out=cv.glmnet(x[train,],y[train],alpha=1)
plot(cv.out)
```



```
bestlam=cv.out$lambda.min
```

```
bestlam
```

```
[1] 25.72381
```

```
ridge.pred=predict(ridge.mod,s=bestlam,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
[1] 1145142
```

Principal Component Regression

```
set.seed(1)
```

```
pcr.fit=pcr(Apps~.,data=College,scale=TRUE,validation="CV")
pcr.fit
```


Principal component regression , fitted with the singular value decomposition algorithm.
Cross-validated using 10 random segments.


Call:

```
pcr(formula = Apps ~ ., data = College, scale = TRUE, validation = "CV")
```

```
# try on training & testing
```

```
pcr.fit=pcr(Apps~.,data=College,subset=train,scale=TRUE,validation="CV")
```

```
validationplot(pcr.fit,val.type="MSEP")
```



```
pcr.pred=predict(pcr.fit,x[test,],ncomp=17)  
mean((pcr.pred-y.test)^2)
```

```
[1] 1166368
```